

# IEMS-469 Dynamic Programming HW2

Weijia Zhao <sup>1</sup>  
Kellogg Finance

This version: November 9, 2021

---

<sup>1</sup>Weijia.Zhao@kellogg.northwestern.edu

Outline: Due to the difficulties to access Deepdish computing resources (too crowded) and the fact that I cannot move files using either Cyberduck/FileZilla (says I am not authorized to do so), I use Northwestern Quest for the the training of question (b) Pong-v0 in this homework assignment. (I have an existing type I account and the GPU used for this homework is A100). Question (a) is fairly easy and the computation is extremely fast even just using CPU.

I use a moving average with a constant decay rate to calculate the average reward versus training epochs rather than only consider the most recent 10 epochs or so as


- (1) Moving average technically takes into account all historical episodes and is less random than the average of 10 most recent epochs (eg: when you just get 10 consecutive good results by luck)
- (2) Moving average can be easily implemented iteratively along with the training process with lower computational effort:  $MA_t = (1 - \alpha)MA_{t-1} + \alpha R_t$

## 1 Cartpole

As mentioned, this model is trained with CPU due to its simplicity. A base also seems unnecessary (also the hw seems only require a base for Q2) since a model without a base can converge to optimal within 500 epochs. I do provide a file with base in "junk/1\_cartpole\_base.py" along with its output "junk/1\_cartpole.txt" in case you are looking for it, apparently the pattern that adding a base slows down the convergence is not universal and it is likely related to the fact that I reduce the learning rate to accommodate much larger loss (likely also larger derivatives) after adding value losses. Nevertheless, I did not try to optimize the hyperparameter in this questions as it converges within 1 or several minutes even in the worst case scenario. Using a base here does not improve the volatility much either since the output without a baselien is already pretty good.

The plot of individual epoch reward is given in the following graph: as expected, policy gradient method is quite volatile. My initial value for the moving average smoother is 10

Figure 1: Cartpole game




(a) Policy Gradient

## 2 Pong

As mentioned, I train the model with Northwestern Quest due to inaccessibility of Deepdish. I use a relatively small learning rate  $1e-5$  after some trials (as I consulted you the last Friday, even a learning rate  $1e-3$  seems to be too large with a convolutional layer and it quickly goes to boundary and loss stays at 0 after that). I stopped the training after 5000 epochs in about 9 hours, somehow this is slower than I originally expected. Also seems to me the program is relatively faster at the beginning and then as the model improves, each epoch takes much longer time. My explanation is that my agent plays badly at the beginning so it quickly lose the game but later the agent's skill improves and it take much longer time to reach a win/loss between my agent and the computer's agent. (Please let me know if there are any other reasons). Overall the model is learning since the average score increases over time. My initial value for the moving average smoother is 10.

The first time a current reward is greater than 0 is episode 1750. The first time a running reward is greater than 0 is episode 2732. Seems to me the improvement of performance slows down after 4000 epochs though it still possible that here is just a saddle point/local optimal and the model will eventually cross it and improve again later.

Figure 2: Pong game



(a) Policy Gradient